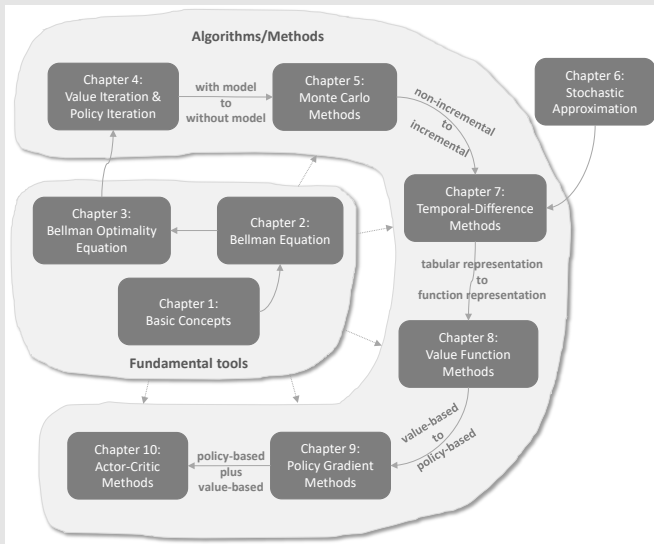# Lecture 7: Temporal-Difference Learning

Shiyu Zhao

Department of Artificial Intelligence
Westlake University

- This lecture introduces temporal-difference (TD) learning, which is one of the most well-known methods in reinforcement learning (RL).

- Monte Carlo (MC) learning is the first model-free method. TD learning is the second model-free method. TD has some advantages compared to MC.

- We will see how the stochastic approximation methods studied in the last lecture are useful.

Outline

Outline

We next consider some stochastic problems and show how to use the RM algorithm to solve them.

First, revisit the mean estimation problem: calculate

$$w = \mathbb{E}[X]$$

based on some iid samples $\{x\}$ of X. We studied it in the last lecture.

- By writing $g(w) = w - \mathbb{E}[X]$, we can reformulate the problem to a root-finding problem

求

$$g(w) = 0$$

- Since we can only obtain samples $\{x\}$ of X, the noisy observation is

观测到的数据 X 有 noise)

$$\tilde{g}(w, \eta) = w - x = (w - \mathbb{E}[X]) + (\mathbb{E}[X] - x) \doteq g(w) + \eta$$

- According to the last lecture, we know the RM algorithm for solving $g(w) = 0$ is

用 RM algorithm, 求 w 值, 使 $g(w^*) = 0$

$$w_{k+1} = w_k - \alpha_k \tilde{g}(w_k, \eta_k) = w_k - \alpha_k(w_k - x_k)$$

Second, consider a little more complex problem. That is to estimate the mean of a function $v(X)$,

$$w = \mathbb{E}[v(X)],$$

就是把观测到的数据包装一下
$x \longrightarrow v(x)$

based on some iid random samples $\{x\}$ of X.

- To solve this problem, we define

$$g(w) = w - \mathbb{E}[v(X)]$$

$$\tilde{g}(w, \eta) = w - v(x) = (w - \mathbb{E}[v(X)]) + (\mathbb{E}[v(X)] - v(x)) \doteq g(w) + \eta.$$

- Then, the problem becomes a root-finding problem: $g(w) = 0$. The corresponding RM algorithm is

$$w_{k+1} = w_k - \alpha_k \tilde{g}(w_k, \eta_k) = w_k - \alpha_k[w_k - v(x_k)]$$

Third, consider an even more complex problem: calculate

$$w = \mathbb{E}[R + \gamma v(X)],$$

再包装一下 $X \longrightarrow R + \gamma v(X)$

where R, X are random variables, γ is a constant, and $v(\cdot)$ is a function.

- Suppose we can obtain samples $\{x\}$ and $\{r\}$ of X and R. We define

$$g(w) = w - \mathbb{E}[R + \gamma v(X)],$$

有两个采测值: X 和 r

$$\tilde{g}(w, \eta) = w - [r + \gamma v(x)]$$
$$= (w - \mathbb{E}[R + \gamma v(X)]) + (\mathbb{E}[R + \gamma v(X)] - [r + \gamma v(x)])$$
$$\doteq g(w) + \eta.$$

- Then, the problem becomes a root-finding problem: $g(w) = 0$. The corresponding RM algorithm is

$$w_{k+1} = w_k - \alpha_k \tilde{g}(w_k, \eta_k) = w_k - \alpha_k[w_k - (r_k + \gamma v(x_k))]$$

This algorithm looks like TD algorithms as shown later.

Quick summary:

- The above three examples become more and more complex.

- They can all be solved by the RM algorithm.

- We will see that the TD algorithms have similar expressions.

Outline

Problem statement:

- Given policy π, the aim is to estimate the state values $\{v_\pi(s)\}_{s \in \mathcal{S}}$ under π.

- Experience samples: $(s_0, r_1, s_1, \ldots, s_t, r_{t+1}, s_{t+1}, \ldots)$ or $\{(s_t, r_{t+1}, s_{t+1})\}_t$
  generated by π.

  由一个policy π 生成
  一组数据 (experience) . 类似于 MC algorithm
  中的 episode.

Important notations:

用 $V(S)$ 近似 $V_\pi(S)$

$$v(s) \longrightarrow v_\pi(s)$$
$$\Downarrow$$
$$v(s_t) \longrightarrow v_\pi(s_t)$$
$$\Downarrow$$
$$v_t(s_t) \longrightarrow v_\pi(s_t)$$

在 episode 的第 t 个时刻 . state 是 S_t . state value 是 $V_t(S_t)$

The TD learning algorithm is

$$v_{t+1}(s_t) = v_t(s_t) - \alpha_t(s_t)\Big[v_t(s_t) - [r_{t+1} + \gamma v_t(s_{t+1})]\Big], \tag{1}$$

$$v_{t+1}(s) = v_t(s), \quad \forall s \neq s_t, \tag{2}$$

⟵ 在t时刻,只有 St的 state value 在更新,其它 state的 值保持不变.

where $t = 0, 1, 2, \ldots$.

Here, $v_t(s_t)$ is the estimated state value of $v_\pi(s_t)$; $\alpha_t(s_t)$ is the learning rate of s_t at time t.

- At time t, only the value of the visited state s_t is updated whereas the values of the unvisited states $s \neq s_t$ remain unchanged.

- The update in (2) will be omitted when the context is clear.

The TD algorithm can be annotated as

$$\underbrace{v_{t+1}(s_t)}_{\text{new estimate}} = \underbrace{v_t(s_t)}_{\text{current estimate}} - \alpha_t(s_t)\Big[\overbrace{v_t(s_t) - \underbrace{[r_{t+1} + \gamma v_t(s_{t+1})]}_{\text{TD target } \bar{v}_t}}^{\text{TD error } \delta_t}\Big], \quad (3)$$

Here,

$$\bar{v}_t \doteq r_{t+1} + \gamma v_t(s_{t+1})$$

is called the TD target. 希望 $v_t(s_t)$ 朝 TD target 接近

$$\delta_t \doteq v_t(s_t) - [r_{t+1} + \gamma v_t(s_{t+1})] = v_t(s_t) - \bar{v}_t$$

is called the TD error.

Observation: The new estimate $v_{t+1}(s_t)$ is a combination of the current estimate $v_t(s_t)$ and the TD error.

First, why is $\bar{v}_t$ called the TD target?

That is because the algorithm drives $v(s_t)$ towards $\bar{v}_t$.

To see that, *Vₜ(Jₜ) 在 t 时刻 发生变化，在 t+1 时刻 变成了 Vₜ₊₁(Jₜ)*

$$v_{t+1}(s_t) = v_t(s_t) - \alpha_t(s_t)\big[v_t(s_t) - \bar{v}_t\big]$$

$$\implies v_{t+1}(s_t) - \bar{v}_t = v_t(s_t) - \bar{v}_t - \alpha_t(s_t)\big[v_t(s_t) - \bar{v}_t\big]$$

$$\implies v_{t+1}(s_t) - \bar{v}_t = [1 - \alpha_t(s_t)][v_t(s_t) - \bar{v}_t]$$

$$\implies |v_{t+1}(s_t) - \bar{v}_t| = |1 - \alpha_t(s_t)||v_t(s_t) - \bar{v}_t|$$

Since $\alpha_t(s_t)$ is a small positive number, we have

$$0 < 1 - \alpha_t(s_t) < 1$$

Therefore,   *t+1 时刻   t 时刻*

$$|v_{t+1}(s_t) - \bar{v}_t| \le |v_t(s_t) - \bar{v}_t|$$

which means $v(s_t)$ is driven towards $\bar{v}_t$!

Second, what is the interpretation of the TD error?

$$\delta_t = v_t(s_t) - [r_{t+1} + \gamma v_t(s_{t+1})]$$

等于 $\delta_t = V_t(s_t) - \widetilde{V_t}$

- It reflects the difference between two time steps.

- It reflects the difference between v_t and v_π. To see that, denote

$$\delta_{\pi,t} \doteq v_\pi(s_t) - [r_{t+1} + \gamma v_\pi(s_{t+1})]$$

Note that

贝尔曼公式

$$\mathbb{E}[\delta_{\pi,t}|S_t = s_t] = v_\pi(s_t) - \mathbb{E}\big[\underline{R_{t+1} + \gamma v_\pi(S_{t+1})}|S_t = s_t\big] = 0.$$

- If $v_t = v_\pi$, then δ_t should be zero (in the expectation sense).
- Hence, if δ_t is not zero, then v_t is not equal to v_π.

- The TD error can be interpreted as innovation, which means new information obtained from the experience (s_t, r_{t+1}, s_{t+1}).

Other properties: 取个 policy ⟶ 求 state value

- The TD algorithm in (3) only estimates the state value of a given policy.
  - It does not estimate the action values.
  - It does not search for optimal policies.
- This algorithm will be extended to estimate action values and then search for optimal policies later in this lecture.
- The TD algorithm in (3) is fundamental for understanding more complex TD algorithms.

Q: What does this TD algorithm do mathematically?

A: It is a model-free algorithm for solving the Bellman equation of a given policy π.

- Chapter 2 has introduced the model-based algorithm for solving the Bellman equation: closed-form solution + iterative algorithm.

First, a new expression of the Bellman equation.

The definition of state value of π is

$$v_\pi(s) = \mathbb{E}\big[R + \gamma G | S = s\big], \quad s \in \mathcal{S} \tag{4}$$

where G is discounted return. Since

$$\mathbb{E}[G|S = s] = \sum_a \pi(a|s) \sum_{s'} p(s'|s, a) v_\pi(s') = \mathbb{E}[v_\pi(S')|S = s],$$

where S' is the next state, we can rewrite (4) as

$$v_\pi(s) = \mathbb{E}\big[R + \gamma v_\pi(S')|S = s\big], \quad s \in \mathcal{S}. \tag{5}$$

贝尔曼 Expectation Equation

Equation (5) is another expression of the Bellman equation. It is sometimes called the Bellman expectation equation, an important tool to design and analyze TD algorithms.

相当于 $g(w)=0$，求 w^*

Second, solve the Bellman equation in (5) **using the RM algorithm.**
In particular, by defining

$$g(v(s)) = v(s) - \mathbb{E}\big[R + \gamma v_\pi(S')|s\big],$$

we can rewrite (5) as

$$g(v(s)) = 0.$$

每次可采样 r 和 s'

Since we can only obtain the samples r and s' of R and S', the noisy observation we have is

$$
\begin{aligned}
\tilde{g}(v(s)) &= v(s) - [r + \gamma v_\pi(s')] \\
&= \underbrace{\Big(v(s) - \mathbb{E}\big[R + \gamma v_\pi(S')|s\big]\Big)}_{g(v(s))} + \underbrace{\Big(\mathbb{E}\big[R + \gamma v_\pi(S')|s\big] - [r + \gamma v_\pi(s')]\Big)}_{\eta}.
\end{aligned}
$$

$$RM \not> g(w) = 0 : \quad w_{k+1} = w_k - \alpha_k \widetilde{g(w_k)}$$

Therefore, the RM algorithm for solving $g(v(s)) = 0$ is

$$v_{k+1}(s) = v_k(s) - \alpha_k \tilde{g}(v_k(s))$$
$$= v_k(s) - \alpha_k \Big(v_k(s) - \big[r_k + \gamma v_\pi(s_k') \big] \Big), \quad k = 1, 2, 3, \ldots \quad (6)$$

where $v_k(s)$ is the estimate of $v_\pi(s)$ at the kth step; r_k, s_k' are the samples of R, S' obtained at the kth step.

The RM algorithm in (6) looks very similar to the TD algorithm. However, there are **two differences**.

- Difference 1: The RM algorithm requires $\{(s, r_k, s_k')\}$ for $k = 1, 2, 3, \ldots$.
  - Modification: $\{(s, r_k, s_k')\}$ is changed to $\{(s_t, r_{t+1}, s_{t+1})\}$ so that the algorithm can utilize the sequential samples in an episode.
- Difference 2: The RM algorithm requires $\underline{v_\pi(s_k')}$. ← 不知道 最优 Policy 是什么
  - Modification: $v_\pi(s_k')$ is replaced by an estimate $v_t(s_{t+1})$. 用当前的 Policy 代替

With the above modifications, the RM algorithms becomes exactly the TD algorithm.

> **Theorem (Convergence of TD Learning)**
>
> *By the TD algorithm (1), $v_t(s)$ converges with probability 1 to $v_\pi(s)$ for all $s \in \mathcal{S}$ as $t \to \infty$ if $\sum_t \alpha_t(s) = \infty$ and $\sum_t \alpha_t^2(s) < \infty$ for all $s \in \mathcal{S}$.*

The proof of the theorem can be found in my book.

Remarks:

- This theorem says the state value can be found by the TD algorithm for a given a policy π.

- $\sum_t \alpha_t(s) = \infty$ and $\sum_t \alpha_t^2(s) < \infty$ must be valid for all $s \in \mathcal{S}$.

  - For condition $\sum_t \alpha_t(s) = \infty$: At time step t,
    ◇ If $s = s_t$, then $\alpha_t(s) > 0$;
    ◇ If $s \neq s_t$, then $\alpha_t(s) = 0$.
    As a result, $\sum_t \alpha_t(s) = \infty$ requires every state must be visited an infinite (or sufficiently many) number of times.

  - For condition $\sum_t \alpha_t^2(s) < \infty$: In practice, the learning rate α is often selected as a small constant. In this case, the condition that $\sum_t \alpha_t^2(s) < \infty$ is invalid anymore. When α is constant, it can still be shown that the algorithm converges in the sense of expectation sense.

While TD learning and MC learning are both model-free, what are the **advantages and disadvantages** of TD learning compared to MC learning?

然后从后往前更新 return 值.

TD/Sarsa learning	MC learning
Online: TD learning is online. It can update the state/action values immediately after receiving a reward.	**Offline:** MC learning is offline. It has to wait until an episode has been completely collected.
Continuing tasks: Since TD learning is online, it can handle both episodic and continuing tasks.	**Episodic tasks:** Since MC learning is offline, it can only handle episodic tasks that has terminate states.

Table: Comparison between TD learning and MC learning.

While TD learning and MC learning are both model-free, what are the **advantages and disadvantages** of TD learning compared to MC learning?

TD/Sarsa learning	MC learning
Bootstrapping: TD bootstraps because the update of a value relies on the previous estimate of this value. Hence, it requires initial guesses.	**Non-bootstrapping:** MC is not bootstrapping, because it can directly estimate state/action values without any initial guess.
Low estimation variance: TD has lower than MC because there are fewer random variables. For instance, Sarsa requires $R_{t+1}, S_{t+1}, A_{t+1}$.	**High estimation variance:** To estimate $q_\pi(s_t, a_t)$, we need samples of $R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$. Suppose the length of each episode is L. There are $|\mathcal{A}|^L$ possible episodes.

Table: Comparison between TD learning and MC learning (continued).

- The TD algorithm introduced in the last section can only estimate state values.

- Next, we introduce, Sarsa, an algorithm that can directly estimate action values.

- We will also see how to use Sarsa to find optimal policies.

First, our aim is to estimate the action values of a given policy π.

Suppose we have some experience $\{(s_t, a_t, r_{t+1}, s_{t+1}, a_{t+1})\}_t$.
We can use the following *Sarsa* algorithm to estimate the action values:

$$q_{t+1}(s_t, a_t) = q_t(s_t, a_t) - \alpha_t(s_t, a_t)\Big[q_t(s_t, a_t) - [r_{t+1} + \gamma q_t(s_{t+1}, a_{t+1})]\Big],$$

$$q_{t+1}(s, a) = q_t(s, a), \quad \forall (s, a) \neq (s_t, a_t),$$

where $t = 0, 1, 2, \ldots$.

- $q_t(s_t, a_t)$ is an estimate of $q_\pi(s_t, a_t)$;
- $\alpha_t(s_t, a_t)$ is the learning rate depending on s_t, a_t.

- **Why is this algorithm called Sarsa?** That is because each step of the algorithm involves $(s_t, a_t, r_{t+1}, s_{t+1}, a_{t+1})$. Sarsa is the abbreviation of state-action-reward-state-action.

- **What is the relationship between Sarsa and the previous TD learning algorithm?** We can obtain Sarsa by replacing the state value estimate $v(s)$ in the TD algorithm with the action value estimate $q(s, a)$. As a result, Sarsa is an action-value version of the TD algorithm.

- **What does the Sarsa algorithm do mathematically?** The expression of Sarsa suggests that it is a stochastic approximation algorithm solving the following equation:

$$q_\pi(s, a) = \mathbb{E}\left[R + \gamma q_\pi(S', A')|s, a\right], \quad \forall s, a.$$

使用 action value，而不是 state value.

This is another expression of the Bellman equation expressed in terms of action values. The proof is given in my book.

Theorem (Convergence of Sarsa learning)

By the Sarsa algorithm, $q_t(s, a)$ converges with probability 1 to the action value $q_\pi(s, a)$ as $t \to \infty$ for all (s, a) if $\sum_t \alpha_t(s, a) = \infty$ and $\sum_t \alpha_t^2(s, a) < \infty$ for all (s, a).

Remarks:

- This theorem says that the action values can be found by Sarsa for a given a policy π.

求 s_t 和 a_t 使: $q_t(s_t, a_t) = r_{t+1} + \gamma q_t(s_{t+1}, a_{t+1})$

The ultimate goal of RL is to find optimal policies.

To do that, we can combine Sarsa with a policy improvement step.

The combined algorithm is also called Sarsa.

Pseudocode: Policy searching by Sarsa

For each episode, do

    Generate a_0 at s_0 following $\pi_0(s_0)$

    If s_t $(t = 0, 1, 2, \dots)$ is not the target state, do 　← 到目标后停止

        Collect an experience sample $(r_{t+1}, s_{t+1}, a_{t+1})$ given (s_t, a_t): generate r_{t+1}, s_{t+1} by interacting with the environment; generate a_{t+1} following $\pi_t(s_{t+1})$.

        Update q-value for (s_t, a_t):　← 用的 RM 算法 $Q_{t+1} = Q_t - \alpha_t \delta_t$

$$q_{t+1}(s_t, a_t) = q_t(s_t, a_t) - \alpha_t(s_t, a_t)\Big[q_t(s_t, a_t) - (r_{t+1} + \gamma q_t(s_{t+1}, a_{t+1}))\Big]$$

目标是 $\delta_t = 0$ 时, Q_t 的值.
δ_t = TD误差

        Update policy for s_t:

$$\pi_{t+1}(a|s_t) = 1 - \frac{\epsilon}{|\mathcal{A}(s_t)|}(|\mathcal{A}(s_t)| - 1) \text{ if } a = \arg\max_a q_{t+1}(s_t, a)$$

$$\pi_{t+1}(a|s_t) = \frac{\epsilon}{|\mathcal{A}(s_t)|} \text{ otherwise}$$

        $s_t \leftarrow s_{t+1}$, $a_t \leftarrow a_{t+1}$

Remarks about this algorithm:

- The policy of s_t is updated immediately after $q(s_t, a_t)$ is updated. This is based on the idea of generalized policy iteration.

- The policy is ϵ-greedy instead of greedy to well balance exploitation and exploration.

Be clear about the core idea and complication:

- The core idea is simple: that is to use an algorithm to solve the Bellman equation of a given policy. 用RM选代算去求 action value .

- The complication emerges when we try to find optimal policies and work efficiently.

Task description:

- The task is to find a good path from a specific starting state to the target state.

  - This task is different from all the previous tasks where we need to find out the optimal policy for every state!
  - Each episode starts from the top-left state and end in the target state.
  - In the future, pay attention to what the task is.

- $r_{\text{target}} = 0$, $r_{\text{forbidden}} = r_{\text{boundary}} = -10$, and $r_{\text{other}} = -1$. The learning rate is $\alpha = 0.1$ and the value of ϵ is 0.1.

Results:

- The left figures above show the final policy obtained by Sarsa.
  - Not all states have the optimal policy.
- The right figures show the total reward and length of every episode.
  - The metric of total reward per episode will be frequently used.

Outline

n-step Sarsa can *unify* Sarsa and Monte Carlo learning

The definition of action value is

$$q_\pi(s, a) = \mathbb{E}[G_t | S_t = s, A_t = a].$$

The discounted return G_t can be written in different forms as

Sarsa $\longleftarrow$ $\quad G_t^{(1)} = R_{t+1} + \gamma q_\pi(S_{t+1}, A_{t+1}),$

$$G_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 q_\pi(S_{t+2}, A_{t+2}),$$

$$\vdots$$

n-step Sarsa $\longleftarrow$ $\quad G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^n q_\pi(S_{t+n}, A_{t+n}),$

在理论上是 等价的，只是不同的分解结构

$$\vdots$$

MC $\longleftarrow$ $\quad G_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \ldots$

It should be noted that $G_t = G_t^{(1)} = G_t^{(2)} = G_t^{(n)} = G_t^{(\infty)}$, where the superscripts merely indicate the different decomposition structures of G_t.

- Sarsa aims to solve

$q_\pi(s,a)$ 是 (s,a) 的 "真实" Return
想让 $q_t(s,a)$ 靠近 $q_\pi(s,a)$

$$q_\pi(s,a) = \mathbb{E}[G_t^{(1)}|s,a] = \mathbb{E}[R_{t+1} + \gamma q_\pi(S_{t+1}, A_{t+1})|s,a].$$

- MC learning aims to solve

$$q_\pi(s,a) = \mathbb{E}[G_t^{(\infty)}|s,a] = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \ldots|s,a].$$

- An intermediate algorithm called n-*step Sarsa* aims to solve

$$q_\pi(s,a) = \mathbb{E}[G_t^{(n)}|s,a] = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^n q_\pi(S_{t+n}, A_{t+n})|s,a].$$

- The algorithm of n-step Sarsa is

$$q_{t+1}(s_t, a_t) = q_t(s_t, a_t)$$
$$- \alpha_t(s_t, a_t)\Big[q_t(s_t, a_t) - [r_{t+1} + \gamma r_{t+2} + \cdots + \gamma^n q_t(s_{t+n}, a_{t+n})]\Big].$$

  - n-step Sarsa becomes the (one-step) Sarsa algorithm when $n = 1$.
  - n-step Sarsa becomes the MC learning algorithm when $n = \infty$.

- Data: n-step Sarsa needs $(s_t, a_t, r_{t+1}, s_{t+1}, a_{t+1}, \ldots, r_{t+n}, s_{t+n}, a_{t+n})$.

- Since $(r_{t+n}, s_{t+n}, a_{t+n})$ has not been collected at time t, we are not able to implement n-step Sarsa at step t. We need to wait until time $t + n$ to update the q-value of (s_t, a_t):

$$q_{t+n}(s_t, a_t) = q_{t+n-1}(s_t, a_t)$$
$$- \alpha_{t+n-1}(s_t, a_t)\Big[q_{t+n-1}(s_t, a_t) - [r_{t+1} + \gamma r_{t+2} + \cdots + \gamma^n q_{t+n-1}(s_{t+n}, a_{t+n})]\Big]$$

- Since n-step Sarsa includes Sarsa and MC learning as two extreme cases, its performance is a blend of Sarsa and MC learning:

  - If n is large, its performance is close to MC learning and hence has a large variance but a small bias.
  - If n is small, its performance is close to Sarsa and hence has a relatively large bias due to the initial guess and relatively low variance.

- Finally, n-step Sarsa is also for policy evaluation. It can be combined with the policy improvement step to search for optimal policies.

Outline

Q-Learning 不需要 Policy Evaluation 和 Policy Improvement

通过迭代 ⟶ 所有(s,a)的 action value ⟶ optimal action value

通过 optimal ⟶ Policy

- Next, we introduce Q-learning, one of the most widely used RL algorithms.
- Sarsa can estimate the action values of a given policy. It must be combined with a policy improvement step to find optimal policies.
- Q-learning can directly estimate optimal action values and hence optimal policies.

更新被访问的 (s,a) 的 q 值。
和 Sarsa 相比，TD target 改变了。

The Q-learning algorithm is

$$q_{t+1}(s_t, a_t) = q_t(s_t, a_t) - \alpha_t(s_t, a_t) \left[q_t(s_t, a_t) - [r_{t+1} + \gamma \max_{a \in \mathcal{A}} q_t(s_{t+1}, a)] \right],$$

$$q_{t+1}(s, a) = q_t(s, a), \quad \forall (s, a) \neq (s_t, a_t), \longleftarrow 不被访问的 (s,a)。其 q 不变。$$

Q-learning is very similar to Sarsa. They are different only in terms of the TD target:

- The TD target in Q-learning is $r_{t+1} + \gamma \max_{a \in \mathcal{A}} q_t(s_{t+1}, a)$
- The TD target in Sarsa is $r_{t+1} + \gamma q_t(s_{t+1}, a_{t+1})$

Sarsa 在解决贝尔曼方程，用的是 a_{t+1}

Q-Learning 在解决贝尔曼最优方程，用的是 能 max q 的 a.

What does Q-learning do mathematically?

It aims to solve

$$q(s, a) = \mathbb{E}\left[R_{t+1} + \gamma \max_a q(S_{t+1}, a) \Big| S_t = s, A_t = a \right], \quad \forall s, a.$$

This is the Bellman optimality equation expressed in terms of action values.
See the proof in my book.

As a result, Q-learning can directly estimate the optimal action values instead of action values of a given policy.

Before further studying Q-learning, we first introduce two important concepts: on-policy learning and off-policy learning.

There exist two policies in a TD learning task: → 和环境交互,生成 experience (eg. episodes)

- The behavior policy is used to generate experience samples.
- The target policy is constantly updated toward an optimal policy.

→ 更新到最优策略

On-policy vs off-policy:

- When the behavior policy is the same as the target policy, such kind of learning is called on-policy. ⟶ behaviour = target policy
- When they are different, the learning is called off-policy.

Advantages of off-policy learning:

- It can search for optimal policies based on the experience samples generated by any other policies.

- Example: The behavior policy is exploratory so that we can generate episodes visiting every state-action pair sufficiently many times.

比如是 ε-greedy 的 , gamma 比较大的.

(a) Exploratory behavior policy   (b) Generated episode

How to judge if a TD algorithm is on-policy or off-policy?

- First, check what math problem the algorithm aims to solve.
- Second, check what experience samples the algorithm requires.

It deserves special attention because it may be confusing to beginners.

- Sarsa aims to evaluate a given policy π by solving

$$q_\pi(s,a) = \mathbb{E}\left[R + \gamma q_\pi(S', A')|s,a\right], \quad \forall s,a.$$

where $R \sim p(R|s,a)$, $S' \sim p(S'|s,a)$, $A' \sim \pi(A'|S')$.

- MC aims to evaluate a given policy π by solving

$$q_\pi(s,a) = \mathbb{E}\left[R_{t+1} + \gamma R_{t+2} + \ldots | S_t = s, A_t = a\right], \quad \forall s,a.$$

where the samples are generated by π.

- Both Sarsa and MC are on-policy.

- π is the behavior policy because we need the experience samples generated by π to estimate the action values of π.

- π is also the target policy because it is updated continuously so that it approaches the optimal policy.

生成 optimal policy，但在 计算时没用到 policy

★ Q-learning is off-policy.

- First, Q-learning aims to solve the <u>Bellman optimality equation</u>

$$q(s,a) = \mathbb{E}\left[R_{t+1} + \gamma \max_a q(S_{t+1}, a) \Big| S_t = s, A_t = a\right], \quad \forall s, a.$$

- Second, the algorithm is

$$q_{t+1}(s_t, a_t) = q_t(s_t, a_t) - \alpha_t(s_t, a_t)\left[q_t(s_t, a_t) - [r_{t+1} + \gamma \max_{a \in \mathcal{A}} q_t(s_{t+1}, a)]\right]$$

which requires $(s_t, a_t, r_{t+1}, s_{t+1})$.

- The behavior policy is the one for generating a_t in s_t. It can be any policy.

behaviour policy : $S_t \xrightarrow{\pi} a_t$ $\pi(s) = a$

target policy : $q \longrightarrow$ optimal $q \longrightarrow$ optimal $a \longrightarrow$ optimal π

Since Q-learning is off-policy, it can be implemented in an either off-policy or on-policy fashion.

Pseudocode: Policy searching by Q-learning (on-policy version)

和 sarsa 差不多，除了算法

For each episode, do

    If s_t $(t = 0, 1, 2, \dots)$ is not the target state, do

        Collect the experience sample (a_t, r_{t+1}, s_{t+1}) given s_t: generate a_t following $\pi_t(s_t)$; generate r_{t+1}, s_{t+1} by interacting with the environment.

        Update q-value for (s_t, a_t):

$$q_{t+1}(s_t, a_t) = q_t(s_t, a_t) - \alpha_t(s_t, a_t)\Big[q_t(s_t, a_t) - (r_{t+1} + \gamma \max_a q_t(s_{t+1}, a))\Big]$$

        Update policy for s_t:

        $\pi_{t+1}(a|s_t) = 1 - \frac{\epsilon}{|\mathcal{A}(s_t)|}(|\mathcal{A}(s_t)| - 1)$ if $a = \arg\max_a q_{t+1}(s_t, a)$

        $\pi_{t+1}(a|s_t) = \frac{\epsilon}{|\mathcal{A}(s_t)|}$ otherwise

---

用的是 ϵ-greedy

See the book for more detailed pseudocode.

2个策略

Pseudocode: Optimal policy search by Q-learning (off-policy version)

Goal: Learn an optimal target policy π_T for all states from the experience samples generated by π_b.

For each episode $\{s_0, a_0, r_1, s_1, a_1, r_2, \dots\}$ generated by π_b, do

    For each step $t = 0, 1, 2, \dots$ of the episode, do

        Update q-value for (s_t, a_t):

$$q_{t+1}(s_t, a_t) = q_t(s_t, a_t) - \alpha_t(s_t, a_t)\Big[q(s_t, a_t) - (r_{t+1} + \gamma \max_a q_t(s_{t+1}, a))\Big]$$

        Update target policy for s_t:

            $\pi_{T,t+1}(a|s_t) = 1$ if $a = \arg\max_a q_{t+1}(s_t, a)$

            $\pi_{T,t+1}(a|s_t) = 0$ otherwise

See the book for more detailed pseudocode.

用的是 greedy .

Task description:

- The task in these examples is to find an optimal policy for all the states.

- The reward setting is $r_{\text{boundary}} = r_{\text{forbidden}} = -1$, and $r_{\text{target}} = 1$. The discount rate is $\gamma = 0.9$. The learning rate is $\alpha = 0.1$.

Ground truth: an optimal policy and the corresponding optimal state values.

(a) Optimal policy

(b) Optimal state value

The behavior policy and the generated experience (10^5 steps):

s和action概率都
为0.2,均匀采样
即 $\varepsilon=1$

(a) Behavior policy

(b) Generated episode

The policy found by off-policy Q-learning:

(a) Estimated policy

(b) State value error

The importance of exploration: episodes of 10^5 steps

If the policy is not sufficiently exploratory, the samples are not good.

(a) Behavior policy $\epsilon = 0.5$ (b) Generated episode   (c) Q-learning result

(a) Behavior policy $\epsilon = 0.1$

(b) Generated episode

(c) Q-learning result

(a) Behavior policy $\epsilon = 0.1$

(b) Generated episode

(c) Q-learning result

All the algorithms we introduced in this lecture can be expressed in a unified expression:

[handwritten: 对 (S_t, a_t) 的 action value 的估计值]

[handwritten: $Q_{t+1} = Q_t - \alpha G(Q_t)$ ，求 Q_t 使 $G = 0$]

[handwritten: 本质是 RM 算法：]

[handwritten: 目的让 q_t 接近 $\bar{q}_t$]

$$q_{t+1}(s_t, a_t) = q_t(s_t, a_t) - \alpha_t(s_t, a_t)[q_t(s_t, a_t) - \bar{q}_t]$$

where $\bar{q}_t$ is the *TD target*.

[handwritten: TD target]

[handwritten: action value 的真误值（但也是用数据估计来的）]

Different TD algorithms have different $\bar{q}_t$.

[handwritten: 因为是 $\bar{q}_t$]

Algorithm	Expression of $\bar{q}_t$
Sarsa	$\bar{q}_t = r_{t+1} + \gamma q_t(s_{t+1}, a_{t+1})$
n-step Sarsa	$\bar{q}_t = r_{t+1} + \gamma r_{t+2} + \cdots + \gamma^n q_t(s_{t+n}, a_{t+n})$
Q-learning	$\bar{q}_t = r_{t+1} + \gamma \max_a q_t(s_{t+1}, a)$
Monte Carlo	$\bar{q}_t = r_{t+1} + \gamma r_{t+2} + \ldots$

Remark: The MC method can also be expressed in this unified expression by setting $\alpha_t(s_t, a_t) = 1$. In particular, the expression is $q_{t+1}(s_t, a_t) = \bar{q}_t$.

略:可以把 state value 换为 action value

All the TD algorithms can be viewed as stochastic approximation algorithms solving the Bellman equation or Bellman optimality equation:

Algorithm	Equation to solve
Sarsa	BE: $q_\pi(s,a) = \mathbb{E}\left[R_{t+1} + \gamma q_\pi(S_{t+1}, A_{t+1})\vert S_t = s, A_t = a\right]$
n-step Sarsa	BE: $q_\pi(s,a) = \mathbb{E}[R_{t+1}+\gamma R_{t+2}+\cdots+\gamma^n q_\pi(s_{t+n}, a_{t+n})\vert S_t = s, A_t = a]$
Q-learning	BOE: $q(s,a) = \mathbb{E}\left[R_{t+1} + \gamma \max_a q(S_{t+1}, a)\vert S_t = s, A_t = a\right]$
Monte Carlo	BE: $q_\pi(s,a) = \mathbb{E}[R_{t+1}+\gamma R_{t+2}+\ldots\vert S_t = s, A_t = a]$

Outline

- Introduced various TD learning algorithms
- Their expressions, math interpretations, implementation, relationship, examples
- Unified point of view